

FirstJobOpt

July 7, 2025

1 Big Data project A.Y. 2024-2025 - First Job

1.1 Members

- Giovanni Antonioni
- Luca Rubboli - 0001083742

```
[1]: import org.apache.spark.sql.SparkSession

val spark = SparkSession.builder
  .appName("First job")
  .getOrCreate()

val sc = spark.sparkContext
```

Intitilizing Scala interpreter ...

Spark Web UI available at <http://10.201.100.248:4042>

SparkContext available as 'sc' (version = 3.5.1, master = local[*], app id = [local-1751647432015](#))

SparkSession available as 'spark'

```
[1]: import org.apache.spark.sql.SparkSession
spark: org.apache.spark.sql.SparkSession =
org.apache.spark.sql.SparkSession@298196fd
sc: org.apache.spark.SparkContext = org.apache.spark.SparkContext@7577dd2d
```

1.1.1 Define useful parameters

- Dataset location
- Iterator (defined like this to overcome different names for same columns in dataset)

```
[2]: val decimals: Int = 4
val minimumYearDataset = 2024
val projectDir: String = "/Users/luca/Desktop/Luca/Università/Magistrale/Corsi/
↳BigData/Drivers"
val datasetDir = "dataset"
val outputDir = "output/firstJobOutputOpt"
```

```

val yellowDatasetDir = s"$datasetDir/yellow_cab"
val greenDatasetDir = s"$datasetDir/green_cab"
val fhvDatasetDir = s"$datasetDir/fhv_cab"
val fhvhvDatasetDir = s"$datasetDir/fhvhv_cab"
val datasetDirMap: Map[String, String] = Map("yellow" -> yellowDatasetDir,
↳ "green" -> greenDatasetDir, "fhv" -> fhvDatasetDir, "fhvhv" ->↳
↳ fhvhvDatasetDir)
val datasetIterator: Map[String, (String, String)] = Map(
  "yellow" -> ("tpep_dropoff_datetime", "tpep_pickup_datetime"),
  "green" -> ("lpep_dropoff_datetime", "lpep_pickup_datetime"),
  // ("fhv", "tpep_dropoff_datetime", "tpep_pickup_datetime"),
  // ("fhvhv", "tpep_dropoff_datetime", "tpep_pickup_datetime"),
)

```

```

[2]: decimals: Int = 4
minimumYearDataset: Int = 2024
projectDir: String =
/Users/luca/Desktop/Luca/Università/Magistrale/Corsi/BigData/Drivers
datasetDir: String = dataset
outputDir: String = output/firstJobOutputOpt
yellowDatasetDir: String = dataset/yellow_cab
greenDatasetDir: String = dataset/green_cab
fhvDatasetDir: String = dataset/fhv_cab
fhvhvDatasetDir: String = dataset/fhvhv_cab
datasetDirMap: Map[String,String] = Map(yellow -> dataset/yellow_cab, green ->
dataset/green_cab, fhv -> dataset/fhv_cab, fhvhv -> dataset/fhvhv_cab)
datasetIterator: Map[String,(String, String)] = Map(yellow ->
(tpep_dropoff_datetime,tpep_pickup_datetime), green ->
(lpep_dropoff_datetime,lpep_pickup_datetime))

```

1.2 Define Columns for analysis

- Columns names
- Time zones for overprice
- Columns used in classification for average price calculation
- Columns which values are used in analysis

```

[3]: val colDurationMinutes: String = "duration_minutes"
val colDurationMinutesBinLabel: String = "duration_minutes_bin_label"
val colYear: String = "year"
val colWeekdaySurcharge: String = "weekday_surcharge"
val colAggregateFee: String = "fees"
val colAggregateFeeBin: String = "agg_fee_bin_label"
val colDistanceBin: String = "distance_bin_label"
val colFareAmount: String = "fare_amount"
val colPricePerDistance: String = "cost_per_distance"

```

```

val colPricePerTime: String = "cost_per_time"
val colAvgPricePerDistance: String = "avg_cost_per_distance"
val colAvgPricePerTime: String = "avg_cost_per_time"
val colPricePerDistanceDiff: String = "cost_per_distance_diff"
val colPricePerDistanceDiffPcg: String = "cost_per_distance_diff_pcg"
val colPricePerTimeDiff: String = "cost_per_time_diff"
val colPricePerTimeDiffPcg: String = "cost_per_time_diff_pcg"
val colPricePerDistanceDiffPcgLabel: String = colPricePerDistanceDiffPcg + "_label"
val colPricePerTimeDiffPcgLabel: String = colPricePerTimeDiffPcg + "_label"

val timeZoneOver: String = "overnight"
val timeZones = Map(timeZoneOver -> (20, 6), "regular" -> (6, 20))
val weekDaySurcharge: Double = 2.5

val colDurationOvernightPcg: String = s"${timeZoneOver}_duration_pcg"

val colToUse: Set[String] = Set(
  "tpep_pickup_datetime",
  "tpep_dropoff_datetime",
  "lpep_pickup_datetime",
  "lpep_dropoff_datetime",
  "passenger_count",
  "trip_distance",
  "ratecodeid",
  "store_and_fwd_flag",
  "payment_type",
  "fare_amount",
  "extra",
  "mta_tax",
  "tip_amount",
  "tolls_amount",
  "improvement_surcharge",
  "total_amount",
  "congestion_surcharge",
  "airport_fee")

val colFees: Set[String] = Set(
  "extra",
  "mta_tax",
  "improvement_surcharge",
  "congestion_surcharge",
  "airport_fee")

val colsForClassification: Seq[String] = Seq(
  "passenger_count",
  "store_and_fwd_flag",

```

```

    "payment_type",
    colAggregateFeeBin,
    colDurationMinutesBinLabel,
    colDistanceBin,
    colYear,
    s"${colDurationOvernightPcg}_label",
    colPricePerDistanceDiffPcgLabel,
    colPricePerTimeDiffPcgLabel
  )

  val colsForValuesAnalysis: Seq[String] = Seq(
    "passenger_count",
    "store_and_fwd_flag",
    "payment_type",
    colAggregateFeeBin,
    colDurationMinutesBinLabel,
    colDistanceBin,
    colYear,
    s"${colDurationOvernightPcg}_label",
  )

```

```

[3]: colDurationMinutes: String = duration_minutes
    colDurationMinutesBinLabel: String = duration_minutes_bin_label
    colYear: String = year
    colWeekdaySurcharge: String = weekday_surcharge
    colAggregateFee: String = fees
    colAggregateFeeBin: String = agg_fee_bin_label
    colDistanceBin: String = distance_bin_label
    colFareAmount: String = fare_amount
    colPricePerDistance: String = cost_per_distance
    colPricePerTime: String = cost_per_time
    colAvgPricePerDistance: String = avg_cost_per_distance
    colAvgPricePerTime: String = avg_cost_per_time
    colPricePerDistanceDiff: String = cost_per_distance_diff
    colPricePerDistanceDiffPcg: String = cost_per_distance_diff_pcg
    colPricePerTimeDiff: String = cost_per_time_diff
    colPricePerTimeDiffPcg: String = cost_per_time_diff_pcg
    colPricePerDistanceDiffPcgLabel: String = ...

```

1.2.1 Define preprocess rules

```

[4]: import java.time.LocalDateTime

    val featureFilters: Map[String, Any => Boolean] = Map(
      "passenger_count" -> {
        case i: Int => i > 0
      }
    )

```

```

    case f: Float => val i = f.toInt; i > 0
    case d: Double => val i = d.toInt; i > 0
    case _ => false
  },
  "trip_distance" -> {
    case i: Int => i > 0
    case i: Float => i > 0
    case i: Double => i > 0
    case _ => false
  },
  "ratecodeid" -> {
    case i: Int => (i >= 1 && i <= 6) || i == 99
    case f: Float => val i = f.toInt; (i >= 1 && i <= 6) || i == 99
    case d: Double => val i = d.toInt; (i >= 1 && i <= 6) || i == 99
    case _ => false
  },
  "store_and_fwd_flag" -> {
    case i: String => i == "Y" || i == "N"
    case _ => false
  },
  "payment_type" -> {
    case i: Int => i >= 1 && i <= 6
    case f: Float => val i = f.toInt; i >= 1 && i <= 6
    case d: Double => val i = d.toInt; i >= 1 && i <= 6
    case _ => false
  },
  "fare_amount" -> {
    case i: Int => i > 0
    case i: Float => i > 0
    case i: Double => i > 0
    case _ => false
  },
  "tolls_amount" -> {
    case i: Int => i >= 0 && i < 200
    case i: Float => i >= 0 && i < 200
    case i: Double => i >= 0 && i < 200
    case _ => false
  }
}

val taxFilter: Any => Boolean = {
  case tax: Int => tax >= 0 && tax < 20
  case tax: Float => tax >= 0 && tax < 20
  case tax: Double => tax >= 0 && tax < 20
  case _ => false
}

```

```

val dateFilter: (Any, Int) => Boolean = {
  case (date: LocalDateTime, minimumYearDataset: Int) => val year: Int = date.
    ↪getYear; year >= minimumYearDataset && year <= LocalDateTime.now().getYear
  case _ => false
}

```

```

[4]: import java.time.LocalDateTime
featureFilters: Map[String,Any => Boolean] = Map(trip_distance ->
$Lambda$2388/0x0000000801090040@48de70ce, tolls_amount ->
$Lambda$2393/0x0000000801093840@f2a6e9a, payment_type ->
$Lambda$2391/0x0000000801092040@5b431913, fare_amount ->
$Lambda$2392/0x0000000801093040@37a8f687, passenger_count ->
$Lambda$2387/0x000000080108f040@ccb8b8bc, store_and_fwd_flag ->
$Lambda$2390/0x0000000801091840@50c94901, ratecodeid ->
$Lambda$2389/0x0000000801090840@1301598a)
taxFilter: Any => Boolean = $Lambda$2394/0x0000000801094840@63ed99a3
dateFilter: (Any, Int) => Boolean = $Lambda$2395/0x0000000801095040@3232c33a

```

1.2.2 Utils functions for rdd

```

[5]: import java.time.temporal.ChronoUnit
import java.time.{DayOfWeek, LocalDate}
import org.apache.spark.rdd.RDD
import org.apache.spark.sql.Row
import scala.math.BigDecimal.RoundingMode

def getDatasetPath(localPath: String): String = {
  "file://" + projectDir + "/" + localPath
}

def binColByStepValue(rdd: RDD[Row], indexOfColToDiscrete: Int, stepValue: Int,
  ↪5): RDD[Row] = {
  rdd
    .map { row =>
      val value: Double = row.get(indexOfColToDiscrete) match {
        case i: Int => i.toDouble
        case d: Double => d
        case l: Long => l.toDouble
        case s: String => try { s.toDouble } catch { case _: Throwable =>
          ↪Double.NaN}
        case _ => Double.NaN
      }

      val rawBin = (value / stepValue).toInt * stepValue
      val binBase = if (value < 0 && value % stepValue == 0) rawBin + stepValue,
      ↪else rawBin
    }
}

```

```

        val label = if (value < 0) { s"${(binBase - stepValue).toInt}|${binBase.
        ↳toInt}" } else { s"${binBase.toInt}|${(binBase + stepValue).toInt}" }

        Row.fromSeq(row.toSeq :+ label)
    }
}

val castForFilter: Any => Any = {
    case s: String => if (s.matches(".*^-?\d+\\.\\d+$")) s.toDouble else if (s.
    ↳matches(".*^-?\d+$")) s.toInt else s.trim
    case d: Double => d
    case i: Int => i
    case l: Long => l.toDouble
    case f: Float => f.toDouble
    case b: Boolean => b
    case null => null
    case other => other.toString.trim
}

val preciseBucketUDF: (Map[String, (Int, Int)], LocalDateTime, LocalDateTime,
↳Int) => Map[String, Double] = { (timeZones: Map[String, (Int, Int)], start:
↳LocalDateTime, end: LocalDateTime, decimals: Int) =>

    val overlap: (LocalDateTime, LocalDateTime, LocalDateTime, LocalDateTime,
↳Int) => Double = { (start1: LocalDateTime, end1: LocalDateTime, start2:
↳LocalDateTime, end2: LocalDateTime, decimals: Int) =>
        val overlapStart = if (start1.isAfter(start2)) start1 else start2
        val overlapEnd = if (end1.isBefore(end2)) end1 else end2
        if (overlapEnd.isAfter(overlapStart)) BigDecimal(ChronoUnit.MILLIS.
        ↳between(overlapStart, overlapEnd) / 60000.0).setScale(decimals, RoundingMode.
        ↳HALF_UP).toDouble else 0.0
    }

    var result = timeZones.keys.map(_ -> 0.0).toMap

    if (!(start == null || end == null)) {

        if (!end.isBefore(start)) {

            var current = start.toLocalDate.atStartOfDay

            while (!current.isAfter(end)) {
                val nextDay = current.plusDays(1)

                timeZones
                    .foreach {
                        case (label, (startHour, endHour)) if startHour > endHour => {

```

```

        val bucketStartBeforeMidnight = current.withHour(startHour).
        ↪withMinute(0).withSecond(0).withNano(0)
        val bucketEndBeforeMidnight = current.withHour(23).withMinute(59).
        ↪withSecond(59)
        val bucketStartAfterMidnight = current.withHour(0).withMinute(0).
        ↪withSecond(0).withNano(0)
        val bucketEndAfterMidnight = current.withHour(endHour).
        ↪withMinute(0).withSecond(0).withNano(0)

        val minutesBeforeMidnight = overlap(start, end, ↪
        ↪bucketStartBeforeMidnight, bucketEndBeforeMidnight, decimals)
        val minutesAfterMidnight = overlap(start, end, ↪
        ↪bucketStartAfterMidnight, bucketEndAfterMidnight, decimals)

        result = result.updated(label, result(label) + ↪
        ↪minutesBeforeMidnight + minutesAfterMidnight)
    }
    case (label, (startHour, endHour)) => {
        val bucketStart = current.withHour(startHour).withMinute(0).
        ↪withSecond(0).withNano(0)
        val bucketEnd = if (endHour == 24) current.plusDays(1).
        ↪withHour(0).withMinute(0).withSecond(0).withNano(0) else current.
        ↪withHour(endHour).withMinute(0).withSecond(0).withNano(0)

        val minutes = overlap(start, end, bucketStart, bucketEnd, ↪
        ↪decimals)

        result = result.updated(label, result(label) + minutes)
    }
}

current = nextDay
}
}
}
result
}

val isUSHolidayOrWeekend: LocalDate => Boolean = { date =>
    val month = date.getMonthValue
    val day = date.getDayOfMonth
    val dayOfWeek = date.getDayOfWeek

    val isIndependenceDay = month == 7 && day == 4
    val isChristmas = month == 12 && day == 25
    val isNewYear = month == 1 && day == 1

```



```

    val isLaborDay = month == 9 && dayOfWeek == DayOfWeek.MONDAY && day <= 7

    val isThanksgiving = month == 11 && dayOfWeek == DayOfWeek.THURSDAY && day >= 22 && day <= 28 && ((day - 1) / 7 + 1 == 4)

    isIndependenceDay || isChristmas || isNewYear || isLaborDay || isThanksgiving || dayOfWeek == DayOfWeek.SATURDAY || dayOfWeek == DayOfWeek.SUNDAY
  }

val selectColumns: (RDD[Row], Seq[String], Set[String]) => RDD[Row] = { (rdd, headers, columnsToKeep) =>
  val keepIndexes = headers.zipWithIndex.collect {
    case (col, idx) if columnsToKeep.contains(col) => idx
  }

  rdd
    .map { row =>
      val selectedValues = keepIndexes.map(row.get)
      Row.fromSeq(selectedValues)
    }
}

val removeColumns: (RDD[Row], Seq[String], Set[String]) => RDD[Row] = { (rdd, headers, columnsToRemove) =>
  val lowerHeaders = headers.map(_.toLowerCase)
  val removeSet = columnsToRemove.map(_.toLowerCase)

  val keepIndexes = lowerHeaders.zipWithIndex.collect {
    case (col, idx) if !removeSet.contains(col) => idx
  }

  rdd
    .map { row =>
      val selectedValues = keepIndexes.map(row.get)
      Row.fromSeq(selectedValues)
    }
}

```

```

[5]: import java.time.temporal.ChronoUnit
import java.time.{DayOfWeek, LocalDate}
import org.apache.spark.rdd.RDD
import org.apache.spark.sql.Row
import scala.math.BigDecimal.RoundingMode
getDatasetPath: (localPath: String)String
binColByStepValue: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row],
indexOfColToDiscrete: Int, stepValue:
Int)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]

```

```

castForFilter: Any => Any = $Lambda$2535/0x00000000801108040@4c29f215
preciseBucketUDF: (Map[String,(Int, Int)], java.time.LocalDateTime,
java.time.LocalDateTime, Int) => Map[String,Double] =
$Lambda$2536/0x00000000801108840@62cd2882
isUSHolidayOrWeekend: java.time.LocalDate => Boolean =
$Lambda$2537/0x00000000801109040@1659deb8
selectColumns: (org.apache.spark.rdd.RDD[org.apache.spark.sql.Row], Seq[St...

```

2 Actual job

1) Select dataset [yellow or green]

```

[6]: val name: String = "yellow"
     val (dropoff, pickup) = datasetIterator(name)

```

```

[6]: name: String = yellow
     dropoff: String = tpep_dropoff_datetime
     pickup: String = tpep_pickup_datetime

```

2) Load dataset

```

[7]: val startTime = System.currentTimeMillis()

     val dataset = spark.read.parquet(getDatasetPath(datasetDirMap(name)))
     var headers: Seq[String] = dataset.columns.map(_.toLowerCase)
     val indexesToUse: Seq[Int] = headers.zipWithIndex.collect {
       case (h, i) if colToUse.contains(h.toLowerCase) => i
     }
     headers = headers.filter(head => colToUse.contains(head.toLowerCase))

```

```

[7]: startTime: Long = 1751647451895
     dataset: org.apache.spark.sql.DataFrame = [VendorID: int, tpep_pickup_datetime:
timestamp_ntz ... 17 more fields]
     headers: Seq[String] = ArraySeq(tpep_pickup_datetime, tpep_dropoff_datetime,
passenger_count, trip_distance, ratecodeid, store_and_fwd_flag, payment_type,
fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge,
total_amount, congestion_surcharge, airport_fee)
     indexesToUse: Seq[Int] = ArraySeq(1, 2, 3, 4, 5, 6, 9, 10, 11, 12, 13, 14, 15,
16, 17, 18)
     headers: Seq[String] = ArraySeq(tpep_pickup_datetime, tpep_dropoff_datetime,
passenger_count, trip_distance, ratecodeid, store_and_fwd_flag, payment_type,
fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge,
total_amount, congestion_surcharge, airport_fee)

```

3) Filter taxes and features based on filter conditions previously defined

```

[8]: import org.apache.spark.sql.DataFrame
import org.apache.spark.storage.StorageLevel
import java.time.format.DateTimeFormatter

def transformRDD(dataset: DataFrame, idxs: Seq[Int], castFunc: Any => Any): RDD[Row] = {
  dataset.rdd.map(row => Row.fromSeq(idxs.map(row.get).map(castFunc)))
}

val rdd = transformRDD(dataset, indexesToUse, castForFilter)

def applyFilters(rdd: RDD[Row], headers: Seq[String], colOfFees: Set[String],
  taxFilter: Any => Boolean, featFilter: Map[String, Any => Boolean],
  dateFilter: (Any, Int) => Boolean, dropoff: String, pickup: String,
  minimumYearDataset: Int): RDD[Row] = {
  rdd
    .filter { row =>
      val formatter = DateTimeFormatter.ofPattern("yyyy-MM-dd'T'HH:mm[:ss]")
      headers.zip(row.toSeq).forall { case (header: String, value) =>
        val taxFilterCondition = if (colOfFees.contains(header.toLowerCase))
          taxFilter(value) else true

        featFilter.get(header.toLowerCase) match {
          case Some(filterFunc) => taxFilterCondition && filterFunc(value)
          case None => if (header.equals(pickup) || header.equals(dropoff)) {
            dateFilter(LocalDateTime.parse(row.getAs[String](headers.indexOf(header)).trim, formatter), minimumYearDataset) && taxFilterCondition
          } else taxFilterCondition
        }
      }
    }
}

val rddFiltered = applyFilters(rdd, headers, colFees, taxFilter,
  featureFilters, dateFilter, dropoff, pickup, minimumYearDataset)

```

```

[8]: import org.apache.spark.sql.DataFrame
import org.apache.spark.storage.StorageLevel
import java.time.format.DateTimeFormatter
transformRDD: (dataset: org.apache.spark.sql.DataFrame, idxs: Seq[Int],
castFunc: Any => Any)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[8] at
map at <console>:46
applyFilters: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row], headers:
Seq[String], colOfFees: Set[String], taxFilter: Any => Boolean, featFilter:
Map[String,Any => Boolean], dateFilter: (Any, Int) => Boolean, dropoff: String,

```

```

pickup: String, minimumYearDataset:
Int)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddFiltered: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[9] at filter at ...

```

4) Add duration and timezones

```

[9]: import java.time.Duration

def addDuration(rdd: RDD[Row], headers: Seq[String], pickup: String, dropoff: String,
  decimals: Int): RDD[Row] = {
  rdd
    .map {row =>
      val formatter = DateTimeFormatter.ofPattern("yyyy-MM-dd'T'HH:mm[:ss]")

      val pickupStr = row.getAs[String](headers.indexOf(pickup)).trim
      val dropoffStr = row.getAs[String](headers.indexOf(dropoff)).trim

      val pickupTS = LocalDateTime.parse(pickupStr, formatter)
      val dropoffTS = LocalDateTime.parse(dropoffStr, formatter)
      val durationMillis = Duration.between(pickupTS, dropoffTS).toMillis
      val durationMinutes = BigDecimal(durationMillis / 60000.0).
        setScale(decimals, RoundingMode.HALF_UP).toDouble

      val pickupYear = pickupTS.getYear

      Row.fromSeq(row.toSeq ++ Seq(durationMinutes, pickupYear))
    }
    .filter {
      row => row.getAs[Double](row.toSeq.length - 2) > 0.0
    }
}

val rddDuration = addDuration(rddFiltered, headers, pickup, dropoff, decimals)
headers = headers ++ Seq(colDurationMinutes, colYear)

val rddDurationBin = binColByStepValue(rddDuration, headers.
  indexOf(colDurationMinutes), 5)
headers = headers :+ colDurationMinutesBinLabel

def addTimeZones(rdd: RDD[Row], headers: Seq[String], timezones: Map[String,
  (Int, Int)], weekDaySurcharge: Double, colDuration: String, pickup: String,
  dropoff: String, decimals: Int, preciseBucketUDF: (Map[String, (Int, Int)],
  LocalDateTime, LocalDateTime, Int) => Map[String, Double],
  isUSHolidayOrWeekendTZ: LocalDate => Boolean): RDD[Row] = {
  rdd

```

```

.map { row =>
    val formatter = DateTimeFormatter.ofPattern("yyyy-MM-dd'T'HH:mm[:ss]")

    val timeZonesDuration: Map[String, Double] = preciseBucketUDF(timezones,
    ↪LocalDateTime.parse(row.getAs[String](headers.indexOf(pickup)).trim,
    ↪formatter), LocalDateTime.parse(row.getAs[String](headers.indexOf(dropoff)).
    ↪trim, formatter), decimals)

    val weekday_surcharge: Double = if (isUSHolidayOrWeekendTZ(LocalDateTime.
    ↪parse(row.getAs[String](headers.indexOf(pickup)).trim, formatter).
    ↪toLocalDate)) 0 else weekdaySurcharge

    val colsToAdd: Seq[Double] = timezones.keys.toSeq.flatMap { tz =>
        val duration = timeZonesDuration.getOrElse(tz, 0.0)
        val totalDuration = row.getAs[Double](headers.indexOf(colDuration))
        Seq(duration, BigDecimal(duration * 100 / totalDuration).
    ↪setScale(decimals, RoundingMode.HALF_UP).toDouble)
    }
    Row.fromSeq((row.toSeq ++ colsToAdd) :+ weekday_surcharge)
}

val rddTimeZones = addTimeZones(rddDurationBin, headers, timeZones,
    ↪weekdaySurcharge, colDurationMinutes, pickup, dropoff, decimals,
    ↪preciseBucketUDF, isUSHolidayOrWeekend)

val headersToAdd: Seq[String] = timeZones.keys.toSeq.flatMap { tz =>
    Seq(tz + "_duration", tz + "_duration_pcg")
} :+ colWeekdaySurcharge

headers = headers ++ headersToAdd

```

```

[9]: import java.time.Duration
addDuration: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row], headers:
Seq[String], pickup: String, dropoff: String, decimals:
Int)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddDuration: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[11] at filter at <console>:68
headers: Seq[String] = ArraySeq(tpep_pickup_datetime, tpep_dropoff_datetime,
passenger_count, trip_distance, ratecodeid, store_and_fwd_flag, payment_type,
fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge,
total_amount, congestion_surcharge, airport_fee, duration_minutes, year,
duration_minutes_bin_label, overnight_duration, overnight_duration_pcg,
regular_duration, regular_duration_pcg, weekday_surcharge)
rddDurationBin: org.apache.spa...

```

```
[10]: val colToRemoveTimeZones = Set(pickup, dropoff, "overnight_duration",
    ↪ "regular_duration", "regular_duration_pcg", "ratecodeid", "tip_amount",
    ↪ "tolls_amount", "total_amount")

val rddTimeZonesOpt = removeColumns(rddTimeZones, headers, colToRemoveTimeZones)
var headersTimeZonesOpt = headers.filterNot(col => colToRemoveTimeZones.
    ↪ contains(col.toLowerCase))
```

```
[10]: colToRemoveTimeZones: scala.collection.immutable.Set[String] =
Set(tpep_dropoff_datetime, regular_duration, overnight_duration,
regular_duration_pcg, tolls_amount, tip_amount, total_amount,
tpep_pickup_datetime, ratecodeid)
rddTimeZonesOpt: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[14] at map at <console>:153
headersTimeZonesOpt: Seq[String] = ArraySeq(passenger_count, trip_distance,
store_and_fwd_flag, payment_type, fare_amount, extra, mta_tax,
improvement_surcharge, congestion_surcharge, airport_fee, duration_minutes,
year, duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge)
```

5) Add Aggregate fees and bins

```
[11]: def addAggregateFees(rdd: RDD[Row], headers: Seq[String], colOfFees:
    ↪ Set[String]): RDD[Row] = {
  rdd
    .map { row =>
      val fees = colOfFees
        .filter(col => headers.contains(col.toLowerCase))
        .map(col => row.getAs[Double](headers.indexOf(col.toLowerCase))).sum

      Row.fromSeq(row.toSeq :+ fees)
    }
}

val rddAggFees = addAggregateFees(rddTimeZonesOpt, headersTimeZonesOpt, colFees)
headersTimeZonesOpt = headersTimeZonesOpt :+ colAggregateFee

val rddAggFeesBin = binColByStepValue(rddAggFees, headersTimeZonesOpt.
    ↪ indexOf(colAggregateFee), 2)
headersTimeZonesOpt = headersTimeZonesOpt :+ colAggregateFeeBin
```

```
[11]: addAggregateFees: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row],
headers: Seq[String], colOfFees:
Set[String])org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddAggFees: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[15] at map at <console>:44
headersTimeZonesOpt: Seq[String] = ArraySeq(passenger_count, trip_distance,
store_and_fwd_flag, payment_type, fare_amount, extra, mta_tax,
```

```
improvement_surcharge, congestion_surcharge, airport_fee, duration_minutes,
year, duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge,
fees, agg_fee_bin_label)
rddAggFeesBin: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[16] at map at <console>:40
headersTimeZonesOpt: Seq[String] = ArraySeq(passenger_count, trip_distance,
store_and_f...
```

```
[12]: val colToRemoveAggFees = colFees ++ Set(colAggregateFee)

val rddAggFeesOpt = removeColumns(rddAggFeesBin, headersTimeZonesOpt,
  ↪ colToRemoveAggFees)
var headersAggFeesOpt = headersTimeZonesOpt.filterNot(col => colToRemoveAggFees.
  ↪ contains(col.toLowerCase))
```

```
[12]: colToRemoveAggFees: scala.collection.immutable.Set[String] =
Set(improvement_surcharge, fees, extra, airport_fee, congestion_surcharge,
mta_tax)
rddAggFeesOpt: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[17] at map at <console>:153
headersAggFeesOpt: Seq[String] = ArraySeq(passenger_count, trip_distance,
store_and_fwd_flag, payment_type, fare_amount, duration_minutes, year,
duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge,
agg_fee_bin_label)
```

6) Add price per mile and minute

```
[13]: def addPricePerDistanceAndTime(rdd: RDD[Row], headers: Seq[String],
  ↪ colFareAmount: String, colDuration: String, colDistance: String): RDD[Row] =
  ↪ {
    rdd
      .map { row =>
        val pricePerTime = Math.round(row.getAs[Double](headers.
  ↪ indexOf(colFareAmount)) / row.getAs[Double](headers.indexOf(colDuration)) *
  ↪ 100) / 100.0
        val pricePerDistance = Math.round(row.getAs[Double](headers.
  ↪ indexOf(colFareAmount)) / row.getAs[Double](headers.indexOf(colDistance)) *
  ↪ 100) / 100.0

        Row.fromSeq(row.toSeq ++ Seq(pricePerTime, pricePerDistance))
      }
  }

val rddPriced = addPricePerDistanceAndTime(rddAggFeesOpt, headersAggFeesOpt,
  ↪ colFareAmount, colDurationMinutes, "trip_distance")
```

```
headersAggFeesOpt = headersAggFeesOpt ++ Seq(colPricePerTime,
↳colPricePerDistance)
```

```
[13]: addPricePerDistanceAndTime: (rdd:
org.apache.spark.rdd.RDD[org.apache.spark.sql.Row], headers: Seq[String],
colFareAmount: String, colDuration: String, colDistance:
String)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddPriced: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[18] at map at <console>:44
headersAggFeesOpt: Seq[String] = ArraySeq(passenger_count, trip_distance,
store_and_fwd_flag, payment_type, fare_amount, duration_minutes, year,
duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge,
agg_fee_bin_label, cost_per_time, cost_per_distance)
```

```
[14]: val colToRemovePricePerDistanceAndTime = Set(colFareAmount, colDurationMinutes)

val rddPricePerDistanceAndTimeOpt = removeColumns(rddPriced, headersAggFeesOpt,
↳colToRemovePricePerDistanceAndTime)
var headersPricePerDistanceAndTimeOpt = headersAggFeesOpt.filterNot(col =>
↳colToRemovePricePerDistanceAndTime.contains(col.toLowerCase))
```

```
[14]: colToRemovePricePerDistanceAndTime: scala.collection.immutable.Set[String] =
Set(fare_amount, duration_minutes)
rddPricePerDistanceAndTimeOpt:
org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[19] at map
at <console>:153
headersPricePerDistanceAndTimeOpt: Seq[String] = ArraySeq(passenger_count,
trip_distance, store_and_fwd_flag, payment_type, year,
duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge,
agg_fee_bin_label, cost_per_time, cost_per_distance)
```

7) Add distance bin and duration in overnight time zone

```
[15]: val rddDistBin = binColByStepValue(rddPricePerDistanceAndTimeOpt,
↳headersPricePerDistanceAndTimeOpt.indexOf("trip_distance"), 5)
headersPricePerDistanceAndTimeOpt = headersPricePerDistanceAndTimeOpt :+
↳colDistanceBin

val rddOvernightBin = binColByStepValue(rddDistBin,
↳headersPricePerDistanceAndTimeOpt.indexOf(colDurationOvernightPcg), 5)
headersPricePerDistanceAndTimeOpt = headersPricePerDistanceAndTimeOpt :+
↳(colDurationOvernightPcg + "_label")
```

```
[15]: rddDistBin: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[20] at map at <console>:40
headersPricePerDistanceAndTimeOpt: Seq[String] = ArraySeq(passenger_count,
```



```
trip_distance, store_and_fwd_flag, payment_type, year,
duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge,
agg_fee_bin_label, cost_per_time, cost_per_distance, distance_bin_label,
overnight_duration_pcg_label)
rddOvernightBin: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[21] at map at <console>:40
headersPricePerDistanceAndTimeOpt: Seq[String] = ArraySeq(passenger_count,
trip_distance, store_and_fwd_flag, payment_type, year,
duration_minutes_bin_label, overnight_duration_pcg, weekday_surcharge,
agg_fee_bin_label, cost_per_time, cost_per_distance, distance_bin_lab...
```

```
[16]: val colToRemoveOvernightBin = Set("trip_distance", colDurationOvernightPcg)

val rddOvernightBinOpt = removeColumns(rddOvernightBin,
  ↳ headersPricePerDistanceAndTimeOpt, colToRemoveOvernightBin)
var headersOvernightBinOpt = headersPricePerDistanceAndTimeOpt.filterNot(col =>
  ↳ colToRemoveOvernightBin.contains(col.toLowerCase))
```

```
[16]: colToRemoveOvernightBin: scala.collection.immutable.Set[String] =
Set(trip_distance, overnight_duration_pcg)
rddOvernightBinOpt: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[22] at map at <console>:153
headersOvernightBinOpt: Seq[String] = ArraySeq(passenger_count,
store_and_fwd_flag, payment_type, year, duration_minutes_bin_label,
weekday_surcharge, agg_fee_bin_label, cost_per_time, cost_per_distance,
distance_bin_label, overnight_duration_pcg_label)
```

8) Add key for average calculation based on columns for classification

```
[17]: import org.apache.spark.HashPartitioner

val numPartitions = spark.sparkContext.defaultParallelism
val partitioner = new HashPartitioner(numPartitions)

val actualHeader = headersOvernightBinOpt
def addKey(rdd: RDD[Row], colsClassification: Seq[String], headers:
  ↳ Seq[String]): RDD[(String, Row)] = {
  rdd
    .map { row =>
      val key = colsClassification.filter(col => headers.contains(col.
        ↳ toLowerCase))
        .map(col => row.get(headers.indexOf(col.toLowerCase)))
        .mkString("_")
      (key, row)
    }
}
```

```
}

val rddWithKey = addKey(rddOvernightBinOpt, colsForClassification,
  ↪actualHeader).partitionBy(partitioner).persist(StorageLevel.MEMORY_ONLY)
```

```
[17]: import org.apache.spark.HashPartitioner
numPartitions: Int = 12
partitioner: org.apache.spark.HashPartitioner =
org.apache.spark.HashPartitioner@c
actualHeader: Seq[String] = ArraySeq(passenger_count, store_and_fwd_flag,
payment_type, year, duration_minutes_bin_label, weekday_surcharge,
agg_fee_bin_label, cost_per_time, cost_per_distance, distance_bin_label,
overnight_duration_pcg_label)
addKey: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row],
colsClassification: Seq[String], headers:
Seq[String])org.apache.spark.rdd.RDD[(String, org.apache.spark.sql.Row)]
rddWithKey: org.apache.spark.rdd.RDD[(String, org.apache.spark.sql.Row)] =
ShuffledRDD[24] at partitionBy at <console>:56
```

9) Calculate prices per distance and time

```
[18]: def calculatePrices(rdd: RDD[(String, Row)], headers: Seq[String],
  ↪colPriceDistance: String, colPriceTime: String): RDD[(String, (Double,
  ↪Double, Long))] = {
  rdd
    .mapValues { row =>
      val costPerDistance = row.getAs[Double](headers.indexOf(colPriceDistance))
      val costPerTime = row.getAs[Double](headers.indexOf(colPriceTime))
      (costPerDistance, costPerTime, 1L)
    }
}

val rddForAvg = calculatePrices(rddWithKey, headersOvernightBinOpt,
  ↪colPricePerDistance, colPricePerTime)
```

```
[18]: calculatePrices: (rdd: org.apache.spark.rdd.RDD[(String,
org.apache.spark.sql.Row)], headers: Seq[String], colPriceDistance: String,
colPriceTime: String)org.apache.spark.rdd.RDD[(String, (Double, Double, Long))]
rddForAvg: org.apache.spark.rdd.RDD[(String, (Double, Double, Long))] =
MapPartitionsRDD[25] at mapValues at <console>:43
```

10) Calculate average prices per distance and time

```
[19]: def calculateAvgPrices(rdd: RDD[(String, (Double, Double, Long))], decimals:
  ↪Int): RDD[(String, (Double, Double))] = {
  rdd
```

```

    .reduceByKey {
      case ((d1, t1, c1), (d2, t2, c2)) => (d1 + d2, t1 + t2, c1 + c2)
    }
    .mapValues {
      case (sumDist, sumTime, count) =>
        val avgDist = BigDecimal(sumDist / count).setScale(decimals, BigDecimal.
        ↪ RoundingMode.HALF_UP).toDouble
        val avgTime = BigDecimal(sumTime / count).setScale(decimals, BigDecimal.
        ↪ RoundingMode.HALF_UP).toDouble
        (avgDist, avgTime)
    }
    .filter {
      case (_, (dist, time)) => dist > 0.0 && time > 0.0
    }
  }

val rddWithAvgPrices = calculateAvgPrices(rddForAvg, decimals)

```

[19]: calculateAvgPrices: (rdd: org.apache.spark.rdd.RDD[(String, (Double, Double, Long))], decimals: Int)org.apache.spark.rdd.RDD[(String, (Double, Double))]
 rddWithAvgPrices: org.apache.spark.rdd.RDD[(String, (Double, Double))] =
 MapPartitionsRDD[28] at filter at <console>:49

11) Join average prices to previous rdd

```

[20]: import org.apache.spark.broadcast.Broadcast

val broadcastAvgPrices: Broadcast[Map[String, (Double, Double)]] = spark.
  ↪ sparkContext.broadcast(rddWithAvgPrices.collectAsMap().toMap)

def applyJoin(rdd: RDD[(String, Row)], broadcastMap: Broadcast[Map[String, (
  ↪ (Double, Double)]]): RDD[Row] = {
  rdd
    .flatMap {
      case (key, originalRow) => broadcastMap.value.get(key)
        .map {
          case (avgCostPerDistance, avgCostPerTime) => Row.fromSeq(originalRow.
          ↪ toSeq ++ Seq(avgCostPerDistance, avgCostPerTime))
        }
    }
  }

val rddJoined = applyJoin(rddWithKey, broadcastAvgPrices)

rddWithKey.unpersist()

```

```
headersOvernightBinOpt = headersOvernightBinOpt ++ Seq(colAvgPricePerDistance,
↳ colAvgPricePerTime)
```

```
[20]: import org.apache.spark.broadcast.Broadcast
broadcastAvgPrices: org.apache.spark.broadcast.Broadcast[Map[String, (Double,
Double)]] = Broadcast(4)
applyJoin: (rdd: org.apache.spark.rdd.RDD[(String, org.apache.spark.sql.Row)],
broadcastMap: org.apache.spark.broadcast.Broadcast[Map[String, (Double,
Double)]]): org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddJoined: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[29] at flatMap at <console>:48
headersOvernightBinOpt: Seq[String] = ArraySeq(passenger_count,
store_and_fwd_flag, payment_type, year, duration_minutes_bin_label,
weekday_surcharge, agg_fee_bin_label, cost_per_time, cost_per_distance,
distance_bin_label, overnight_duration_pcg_label, avg_cost_per_distance,
avg_cost_per_time)
```

```
[21]: val colToRemoveJoin = Set("trip_distance", colDurationOvernightPcg)

val rddJoinOpt = removeColumns(rddJoined, headersOvernightBinOpt,
↳ colToRemovePricePerDistanceAndTime)
var headersJoinOpt = headersOvernightBinOpt.filterNot(col => colToRemoveJoin.
↳ contains(col.toLowerCase))
```

```
[21]: colToRemoveJoin: scala.collection.immutable.Set[String] = Set(trip_distance,
overnight_duration_pcg)
rddJoinOpt: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[30] at map at <console>:153
headersJoinOpt: Seq[String] = ArraySeq(passenger_count, store_and_fwd_flag,
payment_type, year, duration_minutes_bin_label, weekday_surcharge,
agg_fee_bin_label, cost_per_time, cost_per_distance, distance_bin_label,
overnight_duration_pcg_label, avg_cost_per_distance, avg_cost_per_time)
```

12) Add price comparison w.r.t. average price and actual price difference

```
[22]: def addPriceComparison(rdd: RDD[Row], headers: Seq[String], colPriceDistance:
↳ String, colAvgPriceDistance: String, colPriceTime: String, colAvgPriceTime:
↳ String, decimals: Int) = {
  rdd.map { row =>
    val priceColsToAdd: Seq[Double] = Seq((colPriceDistance,
↳ colAvgPriceDistance), (colPriceTime, colAvgPriceTime))
    .flatMap { case (colPrice, colAvgPrice) =>
      val price = row.getAs[Double](headers.indexOf(colPrice))
      val priceAvg = row.getAs[Double](headers.indexOf(colAvgPrice))
      val priceDiff = BigDecimal(price - priceAvg).setScale(decimals,
↳ BigDecimal.RoundingMode.HALF_UP).toDouble
```

```

        val priceDiffPcg = BigDecimal(priceDiff / priceAvg * 100).
        ↪setScale(decimals, BigDecimal.RoundingMode.HALF_UP).toDouble

        Seq(priceDiff, priceDiffPcg)
    }
    Row.fromSeq(row.toSeq ++ priceColsToAdd)
}
}

val rddPriceComparison = addPriceComparison(rddJoinOpt, headersJoinOpt, ↵
    ↪colPricePerDistance, colAvgPricePerDistance, colPricePerTime, ↵
    ↪colAvgPricePerTime, decimals)
headersJoinOpt = headersJoinOpt ++ Seq(colPricePerDistanceDiff, ↵
    ↪colPricePerDistanceDiffPcg, colPricePerTimeDiff, colPricePerTimeDiffPcg)

```

[22]: addPriceComparison: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row], headers: Seq[String], colPriceDistance: String, colAvgPriceDistance: String, colPriceTime: String, colAvgPriceTime: String, decimals: Int)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddPriceComparison: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[31] at map at <console>:50
headersJoinOpt: Seq[String] = ArraySeq(passenger_count, store_and_fwd_flag, payment_type, year, duration_minutes_bin_label, weekday_surcharge, agg_fee_bin_label, cost_per_time, cost_per_distance, distance_bin_label, overnight_duration_pcg_label, avg_cost_per_distance, avg_cost_per_time, cost_per_distance_diff, cost_per_distance_diff_pcg, cost_per_time_diff, cost_per_time_diff_pcg)

13) Bin price difference per time and distance

```

[23]: val rddPriceDistBin = binColByStepValue(rddPriceComparison, headersJoinOpt.
    ↪indexOf(colPricePerDistanceDiffPcg), 5)
val rddPriceDistTimeBin = binColByStepValue(rddPriceDistBin, headersJoinOpt.
    ↪indexOf(colPricePerTimeDiffPcg), 5)

headersJoinOpt = headersJoinOpt ++ Seq(colPricePerDistanceDiffPcgLabel, ↵
    ↪colPricePerTimeDiffPcgLabel)

```

[23]: rddPriceDistBin: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[32] at map at <console>:40
rddPriceDistTimeBin: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[33] at map at <console>:40
headersJoinOpt: Seq[String] = ArraySeq(passenger_count, store_and_fwd_flag, payment_type, year, duration_minutes_bin_label, weekday_surcharge, agg_fee_bin_label, cost_per_time, cost_per_distance, distance_bin_label, overnight_duration_pcg_label, avg_cost_per_distance, avg_cost_per_time,

```
cost_per_distance_diff, cost_per_distance_diff_pcg, cost_per_time_diff,
cost_per_time_diff_pcg, cost_per_distance_diff_pcg_label,
cost_per_time_diff_pcg_label)
```

```
[24]: val colToRemovePriceDistTimeBin = Set(colPricePerDistanceDiff,
      ↪ colPricePerDistanceDiffPcg, colPricePerTimeDiff, colPricePerTimeDiffPcg,
      ↪ colWeekdaySurcharge, colPricePerTime, colPricePerDistance,
      ↪ colAvgPricePerDistance, colAvgPricePerTime)

val rddPriceDistTimeBinOpt = removeColumns(rddPriceDistTimeBin, headersJoinOpt,
      ↪ colToRemovePriceDistTimeBin)
var headersPriceDistTimeBinOpt = headersJoinOpt.filterNot(col =>
      ↪ colToRemovePriceDistTimeBin.contains(col.toLowerCase))
```

```
[24]: colToRemovePriceDistTimeBin: scala.collection.immutable.Set[String] =
Set(avg_cost_per_distance, cost_per_distance, cost_per_distance_diff_pcg,
avg_cost_per_time, cost_per_time_diff, weekday_surcharge,
cost_per_time_diff_pcg, cost_per_distance_diff, cost_per_time)
rddPriceDistTimeBinOpt: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[34] at map at <console>:153
headersPriceDistTimeBinOpt: Seq[String] = ArraySeq(passenger_count,
store_and_fwd_flag, payment_type, year, duration_minutes_bin_label,
agg_fee_bin_label, distance_bin_label, overnight_duration_pcg_label,
cost_per_distance_diff_pcg_label, cost_per_time_diff_pcg_label)
```

14) Reduce to analysis columns only

```
[25]: val headersForAnalysis = headersPriceDistTimeBinOpt.zipWithIndex.filter(head =>
      ↪ colsForClassification.contains(head._1.toLowerCase))

val headersForAnalysisIdxs = headersForAnalysis.map(_._2)
val headersForAnalysisCols = headersForAnalysis.map(_._1)

def reduceToAnalysis(rdd: RDD[Row], idxs: Seq[Int]): RDD[Row] = {
  rdd.map { row => Row.fromSeq(idxs.map(row.get)) }
}

val rddAnalysis = reduceToAnalysis(rddPriceDistTimeBinOpt,
      ↪ headersForAnalysisIdxs)

val totalCount = rddAnalysis.count()
```

```
[25]: headersForAnalysis: Seq[(String, Int)] = ArraySeq((passenger_count,0),
(store_and_fwd_flag,1), (payment_type,2), (year,3),
(duration_minutes_bin_label,4), (agg_fee_bin_label,5), (distance_bin_label,6),
```

```
(overnight_duration_pcg_label,7), (cost_per_distance_diff_pcg_label,8),
(cost_per_time_diff_pcg_label,9))
headersForAnalysisIdxs: Seq[Int] = ArraySeq(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)
headersForAnalysisCols: Seq[String] = ArraySeq(passenger_count,
store_and_fwd_flag, payment_type, year, duration_minutes_bin_label,
agg_fee_bin_label, distance_bin_label, overnight_duration_pcg_label,
cost_per_distance_diff_pcg_label, cost_per_time_diff_pcg_label)
reduceToAnalysis: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row],
idxs: Seq[Int])org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
rddAnaly...
```

15) Group by feature value

```
[26]: def groupByFeatures(rdd: RDD[Row], colForValuesAnalysis: Seq[String],
    ↪ colPriceDistanceDiffPcgLabel: String, colPriceTimeDiffPcgLabel: String,
    ↪ headersAnalysis: Seq[String], decimals: Int, totalCount: Long):
    ↪ Seq[RDD[Row]] = {
    colForValuesAnalysis
    .map { colName =>
        val groupCols = Seq(colPriceDistanceDiffPcgLabel,
    ↪ colPriceTimeDiffPcgLabel):+ colName
        val grouped = rdd
        .map { row =>
            val key = groupCols.map(col => row.get(headersAnalysis.indexOf(col.
    ↪ toLowerCase)))
            (key, 1)
        }
        .reduceByKey(_ + _).map { case (keySeq, count) =>
            val value = keySeq.last.toString
            val costDistLabel = keySeq(0).toString
            val costTimeLabel = keySeq(1).toString
            val pcg = BigDecimal(count.toDouble / totalCount * 100).
    ↪ setScale(decimals, BigDecimal.RoundingMode.HALF_UP).toDouble
            Row.fromSeq(Seq(colName, value, count, pcg, costDistLabel,
    ↪ costTimeLabel))
        }
        grouped
    }
}

val rddFeatures = groupByFeatures(rddAnalysis, colsForValuesAnalysis,
    ↪ colPricePerDistanceDiffPcgLabel, colPricePerTimeDiffPcgLabel,
    ↪ headersForAnalysisCols, decimals, totalCount)
```

```
[26]: groupByFeatures: (rdd: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row],
colForValuesAnalysis: Seq[String], colPriceDistanceDiffPcgLabel: String,
```

```
colPriceTimeDiffPcgLabel: String, headersAnalysis: Seq[String], decimals: Int,
totalCount: Long)Seq[org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]]
rddFeatures: Seq[org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]] =
List(MapPartitionsRDD[38] at map at <console>:53, MapPartitionsRDD[41] at map at
<console>:53, MapPartitionsRDD[44] at map at <console>:53, MapPartitionsRDD[47]
at map at <console>:53, MapPartitionsRDD[50] at map at <console>:53,
MapPartitionsRDD[53] at map at <console>:53, MapPartitionsRDD[56] at map at
<console>:53, MapPartitionsRDD[59] at map at <console>:53)
```

16) Reduce to single rdd and write output

```
[27]: import org.apache.spark.sql.types._

val headersForSchema = Seq(
  StructField("feature", StringType),
  StructField("value", StringType),
  StructField("count", IntegerType),
  StructField("pcg", DoubleType),
  StructField("cost_distance_label", StringType),
  StructField("cost_time_label", StringType)
)

val schema = StructType(headersForSchema)

val dfForAnalysis = spark.createDataFrame(rddFeatures.reduce(_ union _).
  ↪ coalesce(1), schema)

dfForAnalysis.show(1)
val endTime = System.currentTimeMillis()
val durationMs = endTime - startTime

println(s"Job $name-dataset optimized executed in $durationMs ms")

dfForAnalysis.write.mode("overwrite").parquet(getDatasetPath(outputDir + f"/
  ↪ $name"))
```

```
+-----+-----+-----+-----+-----+-----+
|      feature|value|count|pcg|cost_distance_label|cost_time_label|
+-----+-----+-----+-----+-----+-----+
|passenger_count|  2.0|   3|0.0|          [110|115)|          [320|325)|
+-----+-----+-----+-----+-----+-----+

only showing top 1 row
```

Job yellow-dataset optimized executed in 1414331 ms

```
[27]: import org.apache.spark.sql.types._
headersForSchema: Seq[org.apache.spark.sql.types.StructField] =
```



```
List(StructField(feature,StringType,true), StructField(value,StringType,true),
StructField(count,IntegerType,true), StructField(pcg,DoubleType,true),
StructField(cost_distance_label,StringType,true),
StructField(cost_time_label,StringType,true))
schema: org.apache.spark.sql.types.StructType = StructType(StructField(feature,StringType,true),StructField(value,StringType,true),StructField(count,IntegerType,true),StructField(pcg,DoubleType,true),StructField(cost_distance_label,StringType,true),StructField(cost_time_label,StringType,true))
dfForAnalysis: org.apache.spark.sql.DataFrame = [feature: string, value: string
... 4 more fields]
endTime: Long = 1751648866226
durationMs: Long = 1414331
```