

# SecondJobOpt

July 7, 2025

## 1 Big Data project A.Y. 2024-2025

### 1.1 Members

- Giovanni Antonioni
- Luca Rubboli - 0001083742

### 1.2 Second job

```
[57]: import org.apache.spark.sql.SparkSession
```

```
val spark = SparkSession.builder  
  .appName("Second job with RDDs")  
  .getOrCreate()
```

```
val sc = spark.sparkContext
```

```
[57]: import org.apache.spark.sql.SparkSession  
spark: org.apache.spark.sql.SparkSession =  
org.apache.spark.sql.SparkSession@6a007083  
sc: org.apache.spark.SparkContext = org.apache.spark.SparkContext@66755c3
```

### 1.3 Definition of parameters for the job

Here are defined the variables used for the snippet.

```
[58]: val datasetName = "green"  
val datasetFolder = "./dataset"  
val outputDir = s"/output/secondJobRDD/$datasetName"  
val pathToFiles = s"$datasetFolder/$datasetName"  
val weatherData = s"$datasetFolder/weather/weather_data_2017_2024.csv"  
val weatherWmoLookup = s"$datasetFolder/weather/wmo_lookup_codes.csv"
```

```
[58]: datasetName: String = green  
datasetFolder: String = ./dataset  
outputDir: String = /output/secondJobRDD/green  
pathToFiles: String = ./dataset/green
```

```
weatherData: String = ./dataset/weather/weather_data_2017_2024.csv
weatherWmoLookup: String = ./dataset/weather/wmo_lookup_codes.csv
```

## 1.4 Columns for the analysis

```
[59]: import org.apache.spark.sql.types._

val commonFields = List(
  StructField("VendorID", IntegerType),
  StructField("fare_amount", DoubleType),
  StructField("tip_amount", DoubleType),
  StructField("payment_type", LongType),
  StructField("trip_distance", DoubleType),
  StructField("total_amount", DoubleType)
)

val schemaYellow = StructType(
  StructField("tpep_pickup_datetime", TimestampType) ::
  StructField("tpep_dropoff_datetime", TimestampType) ::
  commonFields
)

val schemaGreen = StructType(
  StructField("lpep_pickup_datetime", TimestampType) ::
  StructField("lpep_dropoff_datetime", TimestampType) ::
  commonFields
)

[59]: import org.apache.spark.sql.types._
commonFields: List[org.apache.spark.sql.types.StructField] =
List(StructField(VendorID,IntegerType,true),
StructField(fare_amount,DoubleType,true),
StructField(tip_amount,DoubleType,true),
StructField(payment_type,LongType,true),
StructField(trip_distance,DoubleType,true),
StructField(total_amount,DoubleType,true))
schemaYellow: org.apache.spark.sql.types.StructType = StructType(StructField(tpe
p_pickup_datetime,TimestampType,true),StructField(tpep_dropoff_datetime,Timestam
pType,true),StructField(VendorID,IntegerType,true),StructField(fare_amount,Doubl
eType,true),StructField(tip_amount,DoubleType,true),StructField(payment_type,Lon
gType,true),StructField(trip_distance,DoubleType,true),StructField(total_amount,
DoubleType,true))
schemaGreen: org.apache.sp...
```

## 2 Load Datasets

First we want to load the dataset relative to the taxi data.

```
[60]: val projectDir: String = "/Users/giovanniantonioni/IdeaProjects/Drivers"
      def getDatasetPath(localPath: String): String = {
        "file://" + projectDir + "/" + localPath
      }
```

```
[60]: projectDir: String = /Users/giovanniantonioni/IdeaProjects/Drivers
      getDatasetPath: (localPath: String)String
```

```
[61]: val (schema, pickupCol, dropoffCol) = datasetName match {
      case "yellow" => (schemaYellow, "tpep_pickup_datetime", ↵
        ↵"tpep_dropoff_datetime")
      case _         => (schemaGreen, "lpep_pickup_datetime", ↵
        ↵"lpep_dropoff_datetime")
    }
```

```
val loadedDataset = spark.read
  .schema(schema)
  .option("recursiveFileLookup", "true")
  .parquet(getDatasetPath(pathToFiles))
  .select(
    $"VendorID",
    col(pickupCol).alias("pickup_datetime"),
    col(dropoffCol).alias("dropoff_datetime"),
    $"fare_amount",
    $"tip_amount",
    $"payment_type",
    $"trip_distance",
    $"total_amount"
  )
  .na.drop()
  .dropDuplicates()
  .rdd
```

```
[61]: schema: org.apache.spark.sql.types.StructType = StructType(StructField(lpep_pick
up_datetime,TimestampType,true),StructField(lpep_dropoff_datetime,TimestampType,
true),StructField(VendorID,IntegerType,true),StructField(fare_amount,DoubleType,
true),StructField(tip_amount,DoubleType,true),StructField(payment_type,LongType,
true),StructField(trip_distance,DoubleType,true),StructField(total_amount,Double
Type,true))
pickupCol: String = lpep_pickup_datetime
dropoffCol: String = lpep_dropoff_datetime
loadedDataset: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[279] at rdd at <console>:84
```

### 3 Filtering

```
[62]: import org.apache.spark.sql.Row
import org.apache.spark.rdd.RDD

def filterDataset(dataset: RDD[Row], name: String): RDD[Row] = {
  val allowedYellowVendorId = Set(1, 2, 6, 7)
  val allowedGreenVendorId = Set(1, 2, 6)

  dataset.filter { case row =>
    val allowedIds = if (name == "yellow") allowedYellowVendorId else
    ↪allowedGreenVendorId
    val vendorId = row.getInt(0)
    allowedIds.contains(vendorId)
  }
  .filter(row => row.getDouble(3) > 0)
  .filter(row => row.getDouble(4) >= 0)
  .filter(row => row.getDouble(4) <= row.getDouble(3) * 1.5)
  .filter(row => row.getDouble(6) > 0)
  .filter{ row =>
    val dropOffDateTime = row.getTimestamp(2)
    val pickupDateTime = row.getTimestamp(1)
    dropOffDateTime.after(pickupDateTime)
  }
}

val filtered = filterDataset(loadedDataset, datasetName)
```

```
[62]: import org.apache.spark.sql.Row
import org.apache.spark.rdd.RDD
filterDataset: (dataset: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row],
name: String)org.apache.spark.rdd.RDD[org.apache.spark.sql.Row]
filtered: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
MapPartitionsRDD[285] at filter at <console>:77
```

```
[63]: val withTripDuration = filtered.map { row =>
  val durationMin = (row.getTimestamp(2).getTime - row.getTimestamp(1).getTime).
  ↪toDouble / (1000 * 60)
  (row, durationMin)
}
```

```
[63]: withTripDuration: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, Double)] =
MapPartitionsRDD[286] at map at <console>:61
```

```
[64]: val tripDistances = withTripDuration.map { case (row, _) => row.getLong(5).
      ↪toInt }
      val tripDurations = withTripDuration.map { case (_, duration) => duration }

      val distanceDF = tripDistances.toDF("trip_distance")
      val durationDF = tripDurations.toDF("trip_duration")

      val tripDistanceOutlier = distanceDF.stat.approxQuantile("trip_distance",
      ↪Array(0.02, 0.98), 0.01)
      val tripDurationOutlier = durationDF.stat.approxQuantile("trip_duration",
      ↪Array(0.02, 0.98), 0.01)
```

```
[64]: tripDistances: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[287] at map at
      <console>:65
      tripDurations: org.apache.spark.rdd.RDD[Double] = MapPartitionsRDD[288] at map
      at <console>:66
      distanceDF: org.apache.spark.sql.DataFrame = [trip_distance: int]
      durationDF: org.apache.spark.sql.DataFrame = [trip_duration: double]
      tripDistanceOutlier: Array[Double] = Array(1.0, 2.0)
      tripDurationOutlier: Array[Double] = Array(2.25, 42.083333333333336)
```

```
[74]: def filterOutOutlier(
      dataset: RDD[(Row, Double)],
      tripDistanceOutlier: Array[Double],
      tripDurationOutlier: Array[Double]
    ): RDD[(Row, Double)] = {

      val distanceLower = tripDistanceOutlier(0)
      val distanceUpper = tripDistanceOutlier(1)
      val durationLower = tripDurationOutlier(0)
      val durationUpper = tripDurationOutlier(1)

      dataset.filter { case (row, duration) =>
        val tripDistance = row.getLong(5).toInt
        tripDistance >= distanceLower && tripDistance <= distanceUpper &&
        duration >= durationLower && duration <= durationUpper
      }
    }

      val filteredOut = filterOutOutlier(withTripDuration, tripDistanceOutlier,
      ↪tripDurationOutlier).cache()
      val enriched = filteredOut.map { case (row, duration) =>
        val pickupCalendar = java.util.Calendar.getInstance()
```

```

pickupCalendar.setTime(row.getTimestamp(1))

val hourOfDay = pickupCalendar.get(java.util.Calendar.HOUR_OF_DAY)

val tipAmount = row.getDouble(4)
val totalAmount = row.getDouble(7)

val tipPercentage = if (totalAmount != 0) (tipAmount / totalAmount) * 100
↳ else 0.0

val tripDistance = row.getDouble(6)
val speedMph = if (duration > 0) tripDistance / (duration / 60.0) else 0.0

(row, duration, hourOfDay, tipPercentage, speedMph)
}

```

[74]: filterOutOutlier: (dataset: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, Double)], tripDistanceOutlier: Array[Double], tripDurationOutlier: Array[Double])org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, Double)]  
 filteredOut: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, Double)] = MapPartitionsRDD[366] at filter at <console>:81  
 enriched: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, Double, Int, Double, Double)] = MapPartitionsRDD[367] at map at <console>:90

[66]: val binned = enriched.map { case (row, duration, hourOfDay, tipPercentage,
↳ speedMph) =>
 val binConfigs = Map(
 "trip\_distance" -> (Seq(1.0, 3.0, 6.0), Seq("0-1", "1-3", "3-6", "6+")),
 "trip\_duration\_min" -> (Seq(5.0, 15.0, 30.0), Seq("0-5", "5-15", "15-30",
↳ "30+")),
 "fare\_amount" -> (Seq(5.0, 10.0, 20.0, 40.0), Seq("0-5", "5-10", "10-20",
↳ "20-40", "40+")),
 "tip\_percentage" -> (Seq(5.0, 10.0, 20.0, 30.0), Seq("0-5%", "5-10%",
↳ "10-20%", "20-30%", "30%+")),
 "speed\_mph" -> (Seq(5.0, 15.0, 30.0), Seq("0-5mph", "5-15mph", "15-30mph",
↳ "30mph+"))
 )

 def assignBin(value: Double, bins: Seq[Double], labels: Seq[String]): String
 ↳ = {
 require(labels.length == bins.length + 1, "You need one more label than bin
 ↳ thresholds.")

 if (value < bins.head) labels.head
 else {

```

        val idx = bins.indexWhere(b => value < b)
        if (idx == -1) labels.last
        else labels(idx)
    }
}

val tripDistance = row.getDouble(6)
val tripDistanceBin = assignBin(
    tripDistance,
    binConfigs("trip_distance")._1,
    binConfigs("trip_distance")._2
)

val tripDurationBin = assignBin(
    duration,
    binConfigs("trip_duration_min")._1,
    binConfigs("trip_duration_min")._2
)

val fareAmount = row.getDouble(3)
val fareAmountBin = assignBin(
    fareAmount,
    binConfigs("fare_amount")._1,
    binConfigs("fare_amount")._2
)

val speedBin = assignBin(
    speedMph,
    binConfigs("speed_mph")._1,
    binConfigs("speed_mph")._2
)

def tripHourBucket(hour: Int): String = hour match {
    case h if h >= 0 && h <= 5 => "late_night"
    case h if h >= 6 && h <= 9 => "morning"
    case h if h >= 10 && h <= 15 => "midday"
    case h if h >= 16 && h <= 19 => "evening"
    case _ => "night"
}

val hourBin = tripHourBucket(hourOfDay)
(row, tripDistanceBin, tripDurationBin, fareAmountBin, speedBin, hourBin,
tipPercentage)
}

```

[66]: binned: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, String, String, String, String, String, Double)] = MapPartitionsRDD[309] at map at <console>:61

```
[67]: val weatherFileRDD = spark.read
      .format("CSV")
      .option("header", "true")
      .load(getDatasetPath(weatherData))
      .rdd
      .map { row =>
        val code = row.getString(1).trim.toInt
        val date = row.getString(0).trim
        (code, date)
      }

val wmoLookupFile = spark.read
  .format("CSV")
  .option("header", "true")
  .load(getDatasetPath(weatherWmoLookup))
  .rdd

val wmoLookupPairRDD = wmoLookupFile.map { row =>
  val data = row.getString(0).split(";")
  val code = data(0).trim.toInt
  val description = data(1).trim
  (code, description)
}

import java.time.LocalDate
import java.sql.Timestamp

val transformedWeatherClassRDD = weatherFileRDD
  .join(wmoLookupPairRDD)
  .map(row => {
    val (id, (date, description)) = row
    val formattedDate = LocalDate.parse(date)
    val timestamp = Timestamp.valueOf(formattedDate.atStartOfDay())
    (id, timestamp, description)
  })
```

```
[67]: weatherFileRDD: org.apache.spark.rdd.RDD[(Int, String)] = MapPartitionsRDD[325]
      at map at <console>:72
      wmoLookupFile: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] =
      MapPartitionsRDD[340] at rdd at <console>:81
      wmoLookupPairRDD: org.apache.spark.rdd.RDD[(Int, String)] =
      MapPartitionsRDD[341] at map at <console>:84
      import java.time.LocalDate
      import java.sql.Timestamp
      transformedWeatherClassRDD: org.apache.spark.rdd.RDD[(Int, java.sql.Timestamp,
```



```
String))] = MapPartitionsRDD[345] at map at <console>:96
```

```
[68]: val weatherMap = transformedWeatherClassRDD
      .map { row => (row._2.toLocalDateTime.toLocalDate, row) }
      .collectAsMap()

      val weatherBC = spark.sparkContext.broadcast(weatherMap)

      val rideByDate = binned.map {data =>
        val pickupDateTime = data._1.getTimestamp(1).toLocalDateTime.toLocalDate
        (pickupDateTime, data)
      }
```

```
[68]: weatherMap: scala.collection.Map[java.time.LocalDate,(Int, java.sql.Timestamp,
String)] = Map(2020-06-11 -> (65,2020-06-11 00:00:00.0,Heavy rain), 2018-06-25
-> (3,2018-06-25 00:00:00.0,Clouds generally forming or developing), 2020-11-05
-> (3,2020-11-05 00:00:00.0,Clouds generally forming or developing), 2018-01-04
-> (75,2018-01-04 00:00:00.0,Heavy continuous fall of snowflakes), 2018-11-19 ->
(3,2018-11-19 00:00:00.0,Clouds generally forming or developing), 2023-04-15 ->
(63,2023-04-15 00:00:00.0,Moderate rain), 2024-12-03 -> (1,2024-12-03
00:00:00.0,Clouds generally dissolving or becoming less developed), 2021-04-29
-> (55,2021-04-29 00:00:00.0,Heavy Drizzle), 2018-01-13 -> (63,2018-01-13
00:00:00.0,Moderate rain), 2018-12-27 -> (3,2018-12-27 00:00:00.0,Clouds
generally forming or d...
```

## 4 Join weather and Ride data

```
[69]: import org.apache.spark.broadcast.Broadcast

def joinWithBroadcast(
  dataset: RDD[(Row, String, String, String, String, String, Double)],
  broadcast: Broadcast[scala.collection.Map[LocalDate, (Int, Timestamp,
↳String)]]
) = {
  dataset.map { data =>
    val pickupDateTime = data._1.getTimestamp(1).toLocalDateTime.toLocalDate
    val weatherOpt = broadcast.value.get(pickupDateTime)
    (data, weatherOpt)
  }
}

val joinedWeather = joinWithBroadcast(binned, weatherBC).filter(_._2.isDefined)

val joinedWeatherCleaned = joinedWeather.map {
  case (ride, Some(weather)) => (ride, weather)
}
```

```

}

val finalRDD = joinedWeatherCleaned.map { case (ride, weather) =>
  def generalWeatherLabel(wmoCode: Int): String = wmoCode match {
    case c if Seq(0, 1).contains(c)           => "clear"
    case c if Seq(2, 3, 4).contains(c)        => "cloudy"
    case c if Seq(45, 48).contains(c)         => "foggy"
    case c if (50 to 67).contains(c)          => "rainy"
    case c if (70 to 77).contains(c)          => "snowy"
    case c if (80 to 99).contains(c)          => "stormy"
    case _                                     => "unknown"
  }

  val generalWeather = generalWeatherLabel(weather._1)
  (ride, weather._1, generalWeather)
}

```

[69]: <console>:82: warning: match may not be exhaustive.  
 It would fail on the following input: (\_, None)  
 val joinedWeatherCleaned = joinedWeather.map {  
 ^

```

import org.apache.spark.broadcast.Broadcast
joinWithBroadcast: (dataset: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row,
String, String, String, String, String, Double)], broadcast: org.apache.spark.br
oadcast.Broadcast[scala.collection.Map[java.time.LocalDate,(Int,
java.sql.Timestamp,
String)])org.apache.spark.rdd.RDD[((org.apache.spark.sql.Row, String, String,
String, String, String, Double), Option[(Int, java.sql.Timestamp, String)])]
joinedWeather: org.apache.spark.rdd.RDD[((org.apache.spark.sql.Row, String,
String, String, String, String, Double), Option[(Int, java.sql.Timestamp,
String)])] = MapPartitionsRDD[349] at filter at <console>:80
joinedWeatherCleaned: org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row,
String, String, String, String, String, Double), (Int, java.sql.T...

```

## 5 Export the results

[70]:

```

val binFields = Seq(
  "tripDistanceBin",
  "tripDurationBin",
  "fareAmountBin",
  "speedBin"
)

val keyedRDD = finalRDD.flatMap{ case (ride, code, generalWeather) =>
  binFields.map { field =>

```

```

    val bin = field match {
      case "fareAmountBin" => ride._4
      case "tripDistanceBin" => ride._2
      case "tripDurationBin" => ride._3
      case "speedBin" => ride._5
    }
    val key = s"${field}_$bin"
    (key, ride)
  }
}

```

```

[70]: binFields: Seq[String] = List(tripDistanceBin, tripDurationBin, fareAmountBin,
speedBin)
keyedRDD: org.apache.spark.rdd.RDD[(String, (org.apache.spark.sql.Row, String,
String, String, String, String, Double))] = MapPartitionsRDD[352] at flatMap at
<console>:72

```

```

[71]: import org.apache.spark.HashPartitioner
import org.apache.spark.storage.StorageLevel

val numPartitions = spark.sparkContext.defaultParallelism
val partitioner = new HashPartitioner(numPartitions)

val distributedKeyedRDD = keyedRDD.mapValues{ case ride =>
  val tipPercentage = ride._7
  (tipPercentage, 1L)
}
  .partitionBy(partitioner)
  .persist(StorageLevel.MEMORY_ONLY)

```

```

[71]: import org.apache.spark.HashPartitioner
import org.apache.spark.storage.StorageLevel
numPartitions: Int = 12
partitioner: org.apache.spark.HashPartitioner =
org.apache.spark.HashPartitioner@c
distributedKeyedRDD: org.apache.spark.rdd.RDD[(String, (Double, Long))] =
ShuffledRDD[354] at partitionBy at <console>:77

```

```

[72]: val avgTipRDD = distributedKeyedRDD
  .reduceByKey((a, b) => (a._1 + b._1, a._2 + b._2))
  .map {
    case (id, (sum, count)) => Row(id, sum / count)
  }

distributedKeyedRDD.unpersist()

```

```

val allTipByBinSchema = StructType(Seq(
  StructField("feature", StringType),
  StructField("avg_tip_pct", DoubleType)
))

spark.createDataFrame(avgTipRDD, allTipByBinSchema)
  .write
  .mode("overwrite")
  .parquet(getDatasetPath(s"$outputDir/tip_avg_per_bin/all_features"))

```

[72]: avgTipRDD: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[356] at map at <console>:73  
allTipByBinSchema: org.apache.spark.sql.types.StructType = StructType(StructField(feature,StringType,true),StructField(avg\_tip\_pct,DoubleType,true))

```

[73]: val avgTipByWeather = finalRDD
      .map(r => (r._3, (r._1._7, 1L)))
      .reduceByKey((a, b) => (a._1 + b._1, a._2 + b._2))
      .map { case (weather, (sumTip, count)) => Row(weather, sumTip / count) }

val weatherSchema = StructType(Seq(
  StructField("weather", StringType),
  StructField("avg_tip_pct", DoubleType)
))

spark.createDataFrame(avgTipByWeather, weatherSchema)
  .write
  .mode("overwrite")
  .parquet(getDatasetPath(s"$outputDir/avg_tip_by_weather"))

```

[73]: avgTipByWeather: org.apache.spark.rdd.RDD[org.apache.spark.sql.Row] = MapPartitionsRDD[362] at map at <console>:74  
weatherSchema: org.apache.spark.sql.types.StructType = StructType(StructField(weather,StringType,true),StructField(avg\_tip\_pct,DoubleType,true))